

Machine Learning Introduction Tutorial

COMS3007

Benjamin Rosman / Phumlani Khoza

February 12, 2018

Due date: Monday, 19 February at 2.00pm.

Use your notes and any resources you find online to answer the following questions.

1. State if the following are examples of supervised learning (distinguish between classification or regression), unsupervised learning, or reinforcement learning. Justify your answer.
 - (a) Determining if someone is wearing glasses in an image.
 - (b) Playing a game of chess.
 - (c) Categorising different types of vehicles on the road.
 - (d) Predicting the judgement of a court case.
 - (e) Predicting student performance in a course.
 - (f) Formulating the best treatment plan for a patient.
 - (g) Developing a computer game AI.
 - (h) Predicting tomorrow's BitCoin price.
2. Find and discuss an example of supervised learning being used in a real-world business.
3. Find and discuss an example of unsupervised learning being used in a real-world business.
4. What features would be the most useful in the following learning tasks? Why?
 - (a) Classifying dogs and cats.
 - (b) Classifying images of dogs and cats.
 - (c) Predicting the flight speed of a bird.
 - (d) Predicting student performance in a course.
 - (e) Classifying songs by genre.
 - (f) Predicting which team would win a soccer game.
 - (g) Predicting the number of sales of chocolate in a supermarket.
5. Learning typically requires a lot of training data, and for supervised learning this includes labelled ground truth data (the target values). For each of the problems in (4), how might you go about collecting the ground truth labels for the training data?